

# notebook02-resolvendo\_um\_problema\_de\_regressao

November 17, 2020

## 1 Resolvendo um problema de regressão

### 1.1 Dados do Curso

**Instituição:** IFES

**Curso:** Mestrado Profissional Computação Aplicada

**Professor:** Francisco de Assis Boldt

**Aluno:** Arthur Chisté Lucas

### 1.2 Ambiente

**IDE:** MS Visual Studio Code

**Versão Python:** 3.8.3 64bits com anaconda 2020.07

### 1.3 Introdução

Nesta tarefa, será utilizado um dataset contendo preços de casas, obtido no site Kaggle:

<https://www.kaggle.com/c/house-prices-advanced-regression-techniques/data>

Conforme abaixo, o dataset precisa ser baixado e armazenado no diretório `data/house_prices_dataset`

```
[45]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

url = 'https://github.com/arthurclucas/ReconhecimentoPadroes/blob/main/data/
↪house_prices_dataset/train.csv?raw=true'
dados = pd.read_csv(url)
dados.head()
```

```
[45]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	

3	4	70	RL	60.0	9550	Pave	NaN	IR1
4	5	60	RL	84.0	14260	Pave	NaN	IR1

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

Removendo colunas com poucos dados preenchidos e preenchendo as demais com N/A

```
[46]: dados.fillna(dados.mean(), inplace=True)
dados.fillna('N/A', inplace=True)
#dados.drop('Alley', axis = 1, inplace=True)
#dados.drop('FireplaceQu', axis = 1, inplace=True)
#dados.drop('PoolQC', axis = 1, inplace=True)
#dados.drop('Fence', axis = 1, inplace=True)
#dados.drop('MiscFeature', axis = 1, inplace=True)
dados.columns[dados.isna().any()].tolist()
```

[46]: []

```
[47]: dados.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id              1460 non-null  int64
1   MSSubClass      1460 non-null  int64
2   MSZoning        1460 non-null  object
3   LotFrontage     1460 non-null  float64
4   LotArea         1460 non-null  int64
5   Street          1460 non-null  object
6   Alley           1460 non-null  object
7   LotShape        1460 non-null  object
8   LandContour     1460 non-null  object
```

9	Utilities	1460	non-null	object
10	LotConfig	1460	non-null	object
11	LandSlope	1460	non-null	object
12	Neighborhood	1460	non-null	object
13	Condition1	1460	non-null	object
14	Condition2	1460	non-null	object
15	BldgType	1460	non-null	object
16	HouseStyle	1460	non-null	object
17	OverallQual	1460	non-null	int64
18	OverallCond	1460	non-null	int64
19	YearBuilt	1460	non-null	int64
20	YearRemodAdd	1460	non-null	int64
21	RoofStyle	1460	non-null	object
22	RoofMatl	1460	non-null	object
23	Exterior1st	1460	non-null	object
24	Exterior2nd	1460	non-null	object
25	MasVnrType	1460	non-null	object
26	MasVnrArea	1460	non-null	float64
27	ExterQual	1460	non-null	object
28	ExterCond	1460	non-null	object
29	Foundation	1460	non-null	object
30	BsmtQual	1460	non-null	object
31	BsmtCond	1460	non-null	object
32	BsmtExposure	1460	non-null	object
33	BsmtFinType1	1460	non-null	object
34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1460	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1460	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64

```

57 FireplaceQu      1460 non-null  object
58 GarageType       1460 non-null  object
59 GarageYrBlt      1460 non-null  float64
60 GarageFinish     1460 non-null  object
61 GarageCars       1460 non-null  int64
62 GarageArea       1460 non-null  int64
63 GarageQual       1460 non-null  object
64 GarageCond       1460 non-null  object
65 PavedDrive       1460 non-null  object
66 WoodDeckSF       1460 non-null  int64
67 OpenPorchSF      1460 non-null  int64
68 EnclosedPorch    1460 non-null  int64
69 3SsnPorch        1460 non-null  int64
70 ScreenPorch      1460 non-null  int64
71 PoolArea         1460 non-null  int64
72 PoolQC           1460 non-null  object
73 Fence            1460 non-null  object
74 MiscFeature      1460 non-null  object
75 MiscVal          1460 non-null  int64
76 MoSold           1460 non-null  int64
77 YrSold            1460 non-null  int64
78 SaleType         1460 non-null  object
79 SaleCondition     1460 non-null  object
80 SalePrice        1460 non-null  int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

Posteriormente, transformar as colunas não numéricas, gerando uma coluna para cada valor com 0 - não possui característica, 1 - possui característica

```
[48]: #dados['MSZoning']
      #pd.get_dummies(dados['MSZoning'], prefix='MSZoning')
```

Por hora, separando apenas os dados numéricos para análise

```
[49]: dados = dados.select_dtypes(include=np.number)
      dados
```

```
[49]:
```

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	\
0	1	60	65.0	8450	7	5	
1	2	20	80.0	9600	6	8	
2	3	60	68.0	11250	7	5	
3	4	70	60.0	9550	7	5	
4	5	60	84.0	14260	8	5	
...	...	...	...	...	...	...	
1455	1456	60	62.0	7917	6	5	
1456	1457	20	85.0	13175	6	6	
1457	1458	70	66.0	9042	7	9	

1458	1459	20	68.0	9717	5	6
1459	1460	20	75.0	9937	5	6

	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	...	WoodDeckSF	\
0	2003	2003	196.0	706	...	0	
1	1976	1976	0.0	978	...	298	
2	2001	2002	162.0	486	...	0	
3	1915	1970	0.0	216	...	0	
4	2000	2000	350.0	655	...	192	
...	...	...	...	...	...	...	
1455	1999	2000	0.0	0	...	0	
1456	1978	1988	119.0	790	...	349	
1457	1941	2006	0.0	275	...	0	
1458	1950	1996	0.0	49	...	366	
1459	1965	1965	0.0	830	...	736	



	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	PoolArea	MiscVal	\
0	61	0	0	0	0	0	
1	0	0	0	0	0	0	
2	42	0	0	0	0	0	
3	35	272	0	0	0	0	
4	84	0	0	0	0	0	
...	...	...	...	...	...	...	
1455	40	0	0	0	0	0	
1456	0	0	0	0	0	0	
1457	60	0	0	0	0	2500	
1458	0	112	0	0	0	0	
1459	68	0	0	0	0	0	



	MoSold	YrSold	SalePrice
0	2	2008	208500
1	5	2007	181500
2	9	2008	223500
3	2	2006	140000
4	12	2008	250000
...	...	...	...
1455	8	2007	175000
1456	2	2010	210000
1457	5	2010	266500
1458	4	2010	142125
1459	6	2008	147500



[1460 rows x 38 columns]

Imprime cada uma das características



```
[50]: for i in range(len(dados.columns) - 1):  
      dados.plot.scatter(x=i, y='SalePrice')
```










































































Selecionando apenas algumas características que aparentam realmente influenciar no modelo



```
[51]: dados = dados[['LotArea', 'YearBuilt', 'YearRemodAdd', 'BsmtFinSF1',
    ↳ 'BsmtFinSF2', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
    ↳ 'GrLivArea', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'SalePrice']]



#dados = dados[['GrLivArea', 'SalePrice']]



for i in range(len(dados.columns) - 1):
    dados.plot.scatter(x=i, y='SalePrice')
```






















Treina a base e colhe os resultados

```
[52]: y = dados['SalePrice']
      X = dados.drop('SalePrice', axis = 1)
```

```
modelo = LinearRegression()
modelo.fit(X, y)
ypred = modelo.predict(X)
print(mean_squared_error(y, ypred))
print(modelo.intercept_)
print(modelo.coef_)
```

1657756395.789723

-2029319.208442636

```
[ 3.78041735e-01  4.58446085e+02  5.71251949e+02  1.40963221e+01
 -3.18799322e+00  2.74742801e+01  2.72462693e+01  2.58480117e+01
 -1.15609015e+01  4.15333795e+01  5.60067286e+01  3.05670480e+01
  1.36486018e+01]
```

```
[53]: # Xi = X.iloc[:, 0].to_frame()
      # modelo = LinearRegression()
      # modelo.fit(Xi, y)
      # ypred = modelo.predict(Xi)
      # print(mean_squared_error(y, ypred))
      # print(modelo.intercept_)
```


```
# print(modelo.coef_)
# plt.scatter(Xi, ypred)
# plt.plot(Xi, ypred, 'r')
```

```
[54]: for i in range(len(X.columns) -1):
        Xi = X.iloc[:, i].to_frame()
        modelo = LinearRegression()
        modelo.fit(Xi, y)
        ypred = modelo.predict(Xi)
        print(mean_squared_error(y, ypred))
        print(modelo.intercept_)
        print(modelo.coef_)
        plt.scatter(Xi, ypred)
        plt.plot(Xi, ypred, 'r')
        plt.show()
```

5867752122.509074

158836.1518968766

[2.09997195]




4582376228.725916


-2530308.2457323573

[1375.37346794]






4684989128.859445  
 -3692146.1698673465  
 [1951.29940606]




5365057232.040517


151061.5625256433  
[67.30604049]




6305972098.107117  
181182.02167533064  
[-5.60321424]




3932401923.829737  
63430.62854550623  
[111.10960369]



3991838509.3767366  
36173.4467951213  
[124.50062222]




5663659636.543386  
 160755.8666562706  
 [58.11460255]




6302653388.721356


181165.70963461525  
[-41.83640766]




3139843209.6665273  
18569.02585648728  
[107.13035897]



3855549505.982716  
71357.42140747685  
[231.64561451]



5643036403.904909  
161542.59764040346  
[205.62042374]



```
[55]: def z_function(W0, W1):
    modelo = LinearRegression()
    modelo.fit(X, y)
    Erro = np.empty(W0.shape)
    for j in range(Erro.shape[0]):
        for k in range(Erro.shape[1]):
            ypred = modelo.predict(X)
            mse = mean_squared_error(y, ypred)
            Erro[j][k] = mse
    return Erro

w0 = np.linspace(-1857756395, -1457756395, 30)
w1 = np.linspace(-8, 8, 30)

W0, W1 = np.meshgrid(w0, w1)

Erro = z_function(W0, W1)

plt.ylabel("Erro")
plt.xlabel("W0")
i=8
plt.plot(W0[i,:], Erro[i,:])
plt.show()
```

```
plt.ylabel("Error")
plt.xlabel("W1")
i=8
plt.plot(W1[:,i],Erro[:,i])
plt.show()
```

